

# Indian Institute of Technology - Jodhpur

## Fundamentals of Distributed Systems

### Assignment– 1

#### Part 1 &2

Student Name: Amruj Sai M M

Roll Number: G24AI2072

Github Link:

Part 1 → <https://github.com/Ullivada/Amruj-vector-clock-kv-store/tree/main/vector-clock-kv-store>

Part 2 → <https://github.com/Ullivada/smart-grid-load-balancer.git>

## Tech Stack

- Python 3.11 – Primary programming language for backend logic
- Docker – For containerizing services and simulating distributed nodes
- Docker Compose – For orchestrating multiple services
- Flask – Micro web framework to expose REST APIs
- Requests – To send HTTP messages between services
- Prometheus – For scraping metrics from services
- Grafana – For dashboard visualization of real-time data

**Docker** helps to package and run applications in self-contained units called containers. Think of a container as a complete, isolated package that has everything your application needs: its code, the software it runs on, and all its supporting files. This setup ensures that the application works the same way everywhere, no matter where it deploys, making it consistent and easy to scale.

**Docker Compose** is a tool that simplifies managing applications made of multiple Docker containers. It is using a YAML file to define and run all the services that make up your application together.

## Part 1: Vector Clocks and Causal Consistency

**Objective:** Implement a key-value store across 3 nodes that maintains causal consistency using vector clocks. The system should detect out-of-order messages and delay their processing until causal dependencies are satisfied.

### **Vector Clocks and Causal Ordering**

Vector clocks are data structures that track the causal relationships between events in distributed systems. Each node maintains a vector of counters representing the number of events observed at each node. These are used to detect the order of events and resolve conflicts. Causal ordering ensures that if one event causally affects another, all nodes observe the events in that order. This technique is essential for achieving eventual consistency and ensuring correctness in distributed applications.

## **System Architecture**

- The system consists of three distributed nodes, each implemented as a Python Flask service running in a Docker container.
- Each node maintains its own local key-value store and a vector clock to track causality.
- When a node receives a local PUT request, it updates its clock and broadcasts the change (with its vector clock) to the other nodes.
- Nodes receiving updates compare vector clocks to determine if the message can be applied or should be buffered.
- Buffered messages are checked periodically and applied once their causal dependencies are fulfilled.
- All nodes communicate over a shared Docker network using REST APIs

## **Steps**

1. Created node.py in python. Used Flask to simulate each node in the distributed system.
2. Each node contains:
  - A local key-value store
  - A buffer for delayed messages
  - A vector clock list [i, j, k]
3. The local PUT triggers the node to increment its clock and sends the update + clock to other nodes.
4. Upon receiving a replicated message, the node checks the causal delivery condition. If not met, it buffers the message.
5. Buffered messages are periodically checked for deliverability.
6. Docker Compose launches 3 containers for each node.

# Node.py

```
Project ▾
  > FDS-Assignment1 E:\Amrui - B
  > External Libraries
  ... > Scratches and Consoles

client.py node.py x
1 import os
2 import json
3 import threading
4 import time
5 from typing import Dict, List
6 import requests
7 from flask import Flask, request, jsonify
8
9 class VectorClock(dict): 3 usages
10     def __init__(self, ids: List[int]):
11         super().__init__({i: 0 for i in ids})
12
13     def tick(self, node_id: int): 1 usage
14         self[node_id] += 1
15
16     def update(self, other: "VectorClock"): 1 usage
17         for k in self.keys():
18             self[k] = max(self[k], other.get(k, 0))
19
20     def copy(self): # type: ignore[override]
21         return VectorClock(List(self.keys()))._replace(self)
22
23     # low-level helper used by copy above
24     def _replace(self, mapping): 1 usage
25         for k, v in mapping.items():
26             self[k] = v
27         return self
28
29     # --- causal-delivery predicate -----
30     def is_causally_ready(self, msg_vc: Dict[int, int], sender: int) -> bool: ...
40
41 app = Flask(__name__)
42
43 NODE_ID = int(os.environ["NODE_ID"])
44 PEERS = [p for p in os.environ.get("PEERS", "").split(",") if p]
45 ALL_IDS = [NODE_ID] + [i for i in range(len(PEERS) + 1) if i != NODE_ID]
46 PORT = int(os.environ.get("PORT", 5000 + NODE_ID))
47
48 kv_store: Dict[str, str] = {}
49 clock = VectorClock(ALL_IDS)
50
51 # inbound replication buffer - list of (msg_dict, received_time)
52 buffer: List[Dict] = []
53 lock = threading.Lock()
54
55 @app.route(rule="/put", methods=["POST"])
56 def put():
57     data = request.get_json(force=True)
58     key, value = data["key"], data["value"]
59
60     with lock:
61         # local event = tick own clock
62         clock.tick(NODE_ID)
63         kv_store[key] = value
64         msg = {"sender": NODE_ID, "key": key, "value": value, "vc": dict(clock)}
65
66     _broadcast(endpoint="/replicate", msg)
67     return jsonify({"status": "ok", "vc": msg["vc"]})
68
69 @app.route(rule="/replicate", methods=["POST"])
70 def replicate():
71     msg = request.get_json(force=True)
72     with lock:
73         buffer.append(msg)
74     return "accepted", 202
75
76
77 @app.route(rule="/get", methods=["GET"])
78 def get():
79     key = request.args.get("key")
80     with lock:
81         val = kv_store.get(key)
82         vc = dict(clock)
83     return jsonify({"key": key, "value": val, "vc": vc})
84
85
86
87 def _delivery_worker(): 1 usage
88     while True:
89         time.sleep(0.05)
```

18:52 CRLF UTF-8 4 spaces Python 3.13 (FDS-Assignment1)

Project > FDS-Assignment1 E:\Amruij - ...> External Libraries> Scratches and Consoles

client.py node.py x

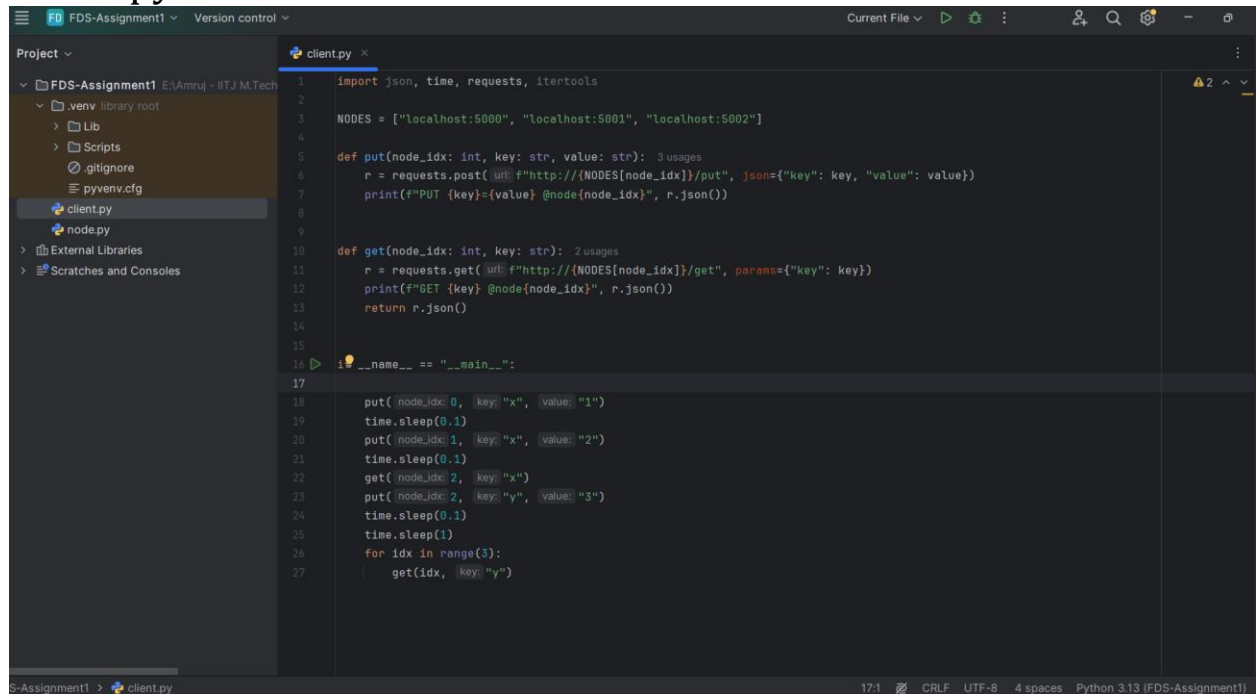
```
1 import os
2 import json
3 import threading
4 import time
5 from typing import Dict, List
6 import requests
7 from flask import Flask, request, jsonify
8
9 class VectorClock(dict): 3 usages
10     def __init__(self, ids: List[int]):
11         super().__init__({i: 0 for i in ids})
12
13     def tick(self, node_id: int): 1 usage
14         self[node_id] += 1
15
16     def update(self, other: "VectorClock"): 1 usage
17         for k in self.keys():
18             self[k] = max(self[k], other.get(k, 0))
19
20     def copy(self): # type: ignore[override]
21         return VectorClock(List(self.keys()))._replace(self)
22
23     # low-level helper used by copy above
24     def _replace(self, mapping): 1 usage
25         for k, v in mapping.items():
26             self[k] = v
27         return self
28
29     # --- causal-delivery predicate -----
30     def is_causally_ready(self, msg_vc: Dict[int, int], sender: int) -> bool: ...
31
32 app = Flask(__name__)
33
34 NODE_ID = int(os.environ["NODE_ID"])
35 PEERS = [p for p in os.environ.get("PEERS", "").split(",") if p]
36 ALL_IDS = [NODE_ID] + [i for i in range(len(PEERS) + 1) if i != NODE_ID]
37 PORT = int(os.environ.get("PORT", 5000 + NODE_ID))
38
39 kv_store: Dict[str, str] = {}
40 clock = VectorClock(ALL_IDS)
41
42 # inbound replication buffer - list of (msg_dict, received_time)
43 buffer: List[Dict] = []
44 lock = threading.Lock()
45
46 @app.route(rule="/put", methods=["POST"])
47 def put():
48     data = request.get_json(force=True)
49     key, value = data["key"], data["value"]
50
51     with lock:
52         # local event => tick own clock
53         clock.tick(NODE_ID)
54         kv_store[key] = value
55         msg = {"sender": NODE_ID, "key": key, "value": value, "vc": dict(clock)}
56
57     _broadcast(endpoint="/replicate", msg)
58     return jsonify({"status": "ok", "vc": msg["vc"]})
59
60 @app.route(rule="/replicate", methods=["POST"])
61 def replicate():
62     msg = request.get_json(force=True)
63     with lock:
64         buffer.append(msg)
65     return "accepted", 202
66
67 @app.route(rule="/get", methods=["GET"])
68 def get():
69     key = request.args.get("key")
70     with lock:
71         val = kv_store.get(key)
72         vc = dict(clock)
73     return jsonify({"key": key, "value": val, "vc": vc})
74
75 def _delivery_worker(): 1 usage
76     while True:
77         time.sleep(0.05)
```

FDS-Assignment1 > node.py 18:52 CRLF UTF-8 4 spaces Python 3.13 (FDS-Assignment1)

node2	node3
 <b>node2</b> • vector-clock-kv-store 5002:5000	 <b>node3</b> • vector-clock-kv-store 5003:5000
 <b>node3</b> • vector-clock-kv-store 5003:5000	 <b>node2</b> • vector-clock-kv-store 5002:5000
 <b>node1</b> • vector-clock-kv-store 5001:5000	 <b>node3</b> • vector-clock-kv-store 5003:5000
	<p>* Serving Flask app 'node' * Debug mode: off</p> <p>WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead. * Running on all addresses (0.0.0.0) * Running on <a href="http://127.0.0.1:5000">http://127.0.0.1:5000</a> * Running on <a href="http://172.22.0.2:5000">http://172.22.0.2:5000</a> Press CTRL+C to quit</p>
	<p>WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead. * Running on all addresses (0.0.0.0) * Running on <a href="http://127.0.0.1:5000">http://127.0.0.1:5000</a> * Running on <a href="http://172.22.0.3:5000">http://172.22.0.3:5000</a> Press CTRL+C to quit</p>
	<p>WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead. * Running on all addresses (0.0.0.0) * Running on <a href="http://127.0.0.1:5000">http://127.0.0.1:5000</a> * Running on <a href="http://172.22.0.4:5000">http://172.22.0.4:5000</a> Press CTRL+C to quit</p>

```
node2-1 | * Serving Flask app 'node'
node2-1 | * Debug mode: off
node3-1 | * Serving Flask app 'node'
node1-1 | * Serving Flask app 'node'
node2-1 | WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
node3-1 | * Debug mode: off
node1-1 | * Debug mode: off
node2-1 | * Running on all addresses (0.0.0.0)
node3-1 | WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
node1-1 | WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
node2-1 | * Running on http://127.0.0.1:5000
node3-1 | * Running on all addresses (0.0.0.0)
node1-1 | * Running on all addresses (0.0.0.0)
node2-1 | * Running on http://172.22.0.4:5000
node3-1 | * Running on http://127.0.0.1:5000
node1-1 | * Running on http://127.0.0.1:5000
node2-1 | Press CTRL+C to quit
node3-1 | * Running on http://172.22.0.2:5000
node1-1 | * Running on http://172.22.0.3:5000
node3-1 | Press CTRL+C to quit
node1-1 | Press CTRL+C to quit
```

# Client.py



```
1 import json, time, requests, itertools
2
3 NODES = ["localhost:5000", "localhost:5001", "localhost:5002"]
4
5 def put(node_idx: int, key: str, value: str): 3 usages
6     r = requests.post(url=f"http://{NODES[node_idx]}/put", json={"key": key, "value": value})
7     print(f"PUT {key}={value} @node{node_idx}", r.json())
8
9
10 def get(node_idx: int, key: str): 2 usages
11     r = requests.get(url=f"http://{NODES[node_idx]}/get", params={"key": key})
12     print(f"GET {key} @node{node_idx}", r.json())
13     return r.json()
14
15
16 if __name__ == "__main__":
17
18     put(node_idx=0, key="x", value="1")
19     time.sleep(0.1)
20     put(node_idx=1, key="x", value="2")
21     time.sleep(0.1)
22     get(node_idx=2, key="x")
23     put(node_idx=2, key="y", value="3")
24     time.sleep(0.1)
25     time.sleep(1)
26     for idx in range(3):
27         get(idx, key="y")
```

E:\Amruj - IITJ M.Tech\T2\FDS\vector-clock-kv-store

## Part 2: Smart Grid Load Balancer

### Summary

This project implements a dynamic load balancing system for a Smart Grid that efficiently distributes Electric Vehicle (EV) charging requests across multiple substations. The system uses real-time load monitoring, intelligent request routing, and comprehensive observability to ensure optimal resource utilization and grid stability. The implemented solution successfully demonstrates:

- **Dynamic Load Balancing:** Intelligent routing based on real-time substation loads
- **Scalable Architecture:** Microservices-based design with containerized deployment
- **Comprehensive Monitoring:** Full observability stack with Prometheus and Grafana
- **High Availability:** Fault-tolerant design with graceful degradation

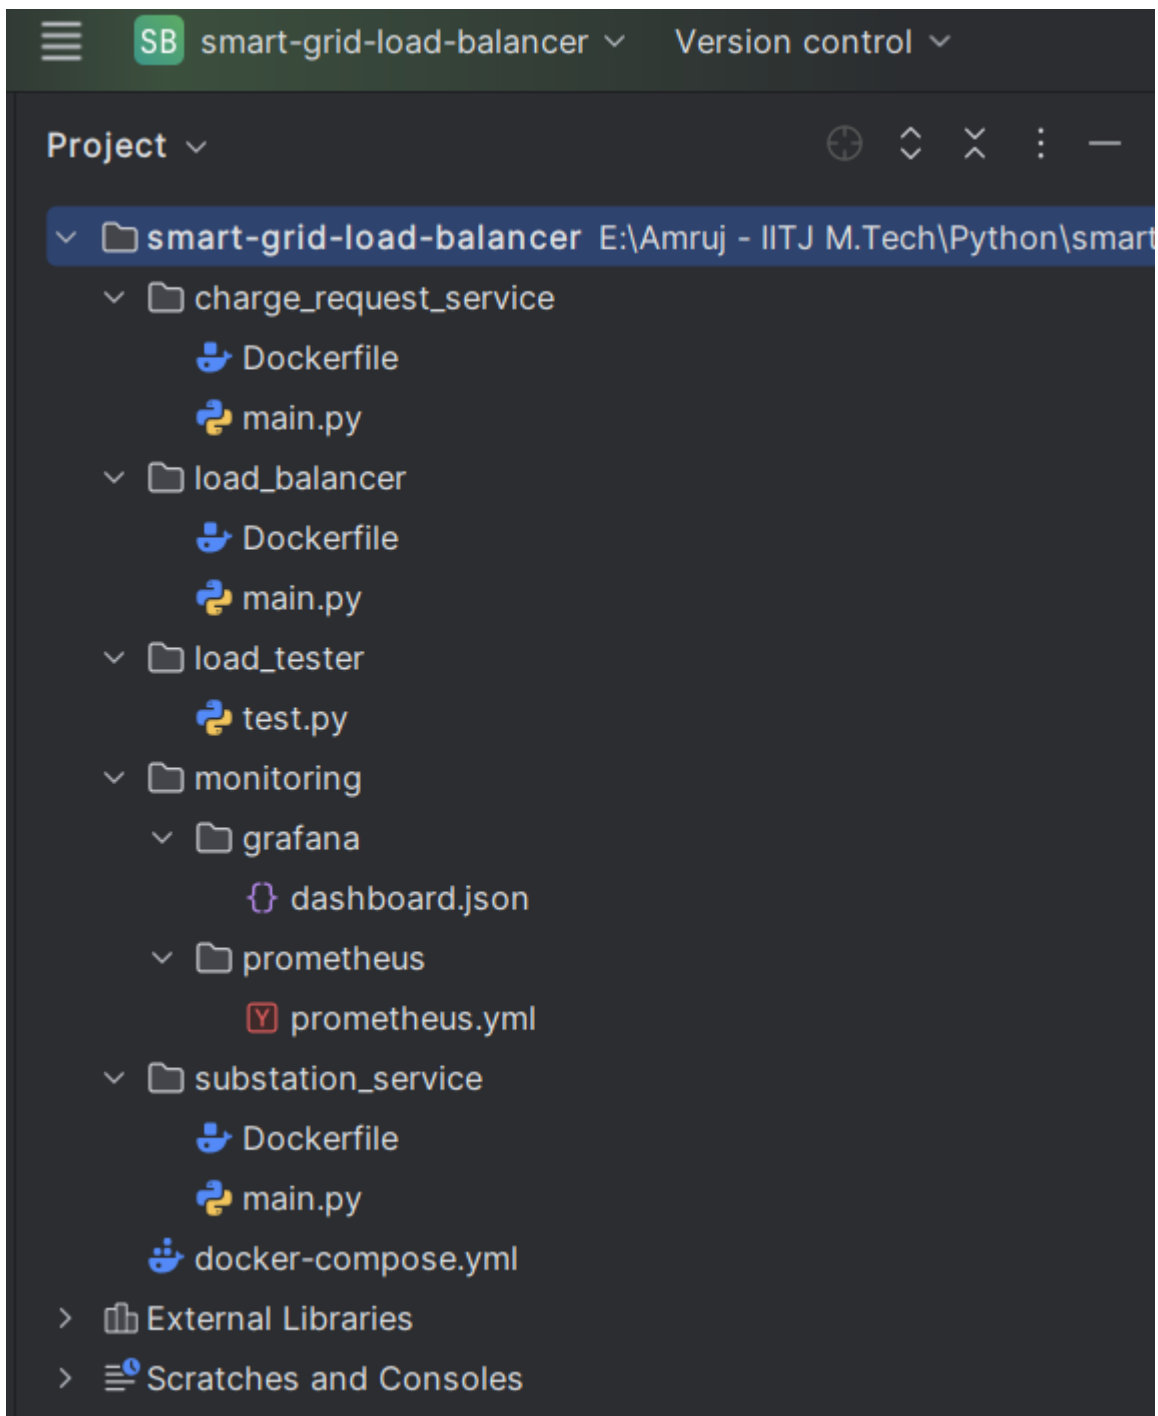
### System Architecture

- The system consists of two substation services, a centralized load balancer, and monitoring tools (Prometheus and Grafana).
- Substations accept EV charging requests and expose their current load via a /metrics endpoint.
- The load balancer periodically polls these metrics and routes new requests to the least-loaded substation.
- Prometheus scrapes the metrics from substations and stores time-series data.
- Grafana connects to Prometheus and visualizes load data through a custom dashboard.

## **Steps**

1. Implemented substation\_service/main.py with /charge and /metrics endpoints.
2. Implemented load\_balancer/main.py to query all substation metrics and route based on lowest load.
3. Created test.py to simulate 50 vehicle requests.
4. Configured Prometheus to scrape metrics from substations.
5. Connected Grafana to Prometheus and imported dashboard.





## charge request service/main.py

```
1 from flask import Flask, request, jsonify
2 import requests
3 import logging
4 import os
5 from datetime import datetime
6
7 app = Flask(__name__)
8 logging.basicConfig(level=logging.INFO)
9 logger = logging.getLogger(__name__)
10
11 LOAD_BALANCER_URL = os.getenv('LOAD_BALANCER_URL', 'http://load_balancer:8080')
12
13 @app.route(rule='/charge', methods=['POST'])
14 def charge_request():
15     """
16     Handles incoming EV charging requests and forwards them to the load balancer
17     """
18     try:
19         # Extract charging request data
20         data = request.get_json()
21
22         if not data:
23             return jsonify({'error': 'No data provided'}), 400
24
25         # Validate required fields
26         required_fields = ['vehicle_id', 'charge_amount', 'priority']
27         for field in required_fields:
28             if field not in data:
29                 return jsonify({'error': f'Missing required field: {field}'}), 400
30
31         # Add timestamp
32         data['timestamp'] = datetime.now().isoformat()
33
34         logger.info(f'Received charge request for vehicle {data["vehicle_id"]}')
35
36         # Forward request to load balancer
37         response = requests.post(
38             url=f'{LOAD_BALANCER_URL}/route_charge',
39             json=data,
40             timeout=30
41         )
42
43         if response.status_code == 200:
44             result = response.json()
45             logger.info(f'Charge request routed successfully to {result.get("substation_id")}')
46             return jsonify(result), 200
47         else:
48             logger.error(f'Load balancer error: {response.status_code}')
49             return jsonify({'error': 'Failed to route charge request'}), 500
50
51     except requests.RequestException as e:
52         logger.error(f'Connection error to load balancer: {str(e)}')
53         return jsonify({'error': 'Load balancer unavailable'}), 503
54     except Exception as e:
55         logger.error(f'Unexpected error: {str(e)}')
56         return jsonify({'error': 'Internal server error'}), 500
57
58 @app.route(rule='/health', methods=['GET'])
59 def health_check():
60     """Health check endpoint"""
61     return jsonify({'status': 'healthy', 'service': 'charge_request_service'}), 200
62
63 @app.route(rule='/status', methods=['GET'])
64 def status():
65     """Service status endpoint"""
66     try:
67         # Check if load balancer is reachable
68         response = requests.get(url=f'{LOAD_BALANCER_URL}/health', timeout=5)
69         lb_status = "healthy" if response.status_code == 200 else "unhealthy"
70     except:
71         lb_status = "unreachable"
72
73     return jsonify({
74         'service': 'charge_request_service',
75         'status': 'running',
76         'load_balancer_status': lb_status,
77         'timestamp': datetime.now().isoformat()
78     }), 200
79
80 if __name__ == '__main__':
```

## load\_balancer/main.py

```
charge_request_service\main.py  test.py  load_balancer\main.py ×
1  from flask import Flask, request, jsonify
2  import requests
3  import logging
4  import os
5  import time
6  import threading
7  from datetime import datetime
8  import re
9
10 app = Flask(__name__)
11 logging.basicConfig(level=logging.INFO)
12 logger = logging.getLogger(__name__)
13
14 # Configuration
15 SUBSTATIONS = [
16     {'id': 'substation_1', 'url': 'http://substation_1:8001'},
17     {'id': 'substation_2', 'url': 'http://substation_2:8001'},
18     {'id': 'substation_3', 'url': 'http://substation_3:8001'}
19 ]
20
21 # Store current loads for each substation
22 substation_loads = {}
23 for substation in SUBSTATIONS:
24     substation_loads[substation['id']] = 0
25
26 # Lock for thread-safe access to load data
27 load_lock = threading.Lock()
28
29 def parse_prometheus_metrics(metrics_text): 1 usage
30     """Parse Prometheus metrics text format"""
31     current_load = 0
32     for line in metrics_text.split('\n'):
33         if line.startswith('substation_current_load'):
34             # Extract the value from the metric line
35             match = re.search(pattern=r'substation_current_load\s+(\d+(?:\.\d+)?)', line)
36             if match:
37                 current_load = float(match.group(1))
38                 break
39     return current_load
40
41 def update_substation_loads(): 1 usage
42     """Periodically update substation loads by polling their metrics"""
mart-grid-load-balancer > load_balancer > main.py
```

```
40
41 def update_substation_loads(): 1 usage
42     """Periodically update substation loads by polling their metrics"""
43     while True:
44         try:
45             for substation in SUBSTATIONS:
46                 try:
47                     response = requests.get(url=f"{substation['url']}/metrics", timeout=5)
48                     if response.status_code == 200:
49                         current_load = parse_prometheus_metrics(response.text)
50                         with load_lock:
51                             substation_loads[substation['id']] = current_load
52                             logger.debug(f"Updated {substation['id']} load: {current_load}")
53                         else:
54                             logger.warning(f"Failed to get metrics from {substation['id']}")
55                     except requests.RequestException as e:
56                         logger.error(f"Error polling {substation['id']}: {str(e)}")
57                         # Keep the last known load value
58
59                 # Wait before next polling cycle
60                 time.sleep(5) # Poll every 5 seconds
61
62         except Exception as e:
63             logger.error(f"Error in load update thread: {str(e)}")
64             time.sleep(10) # Wait longer on error
65
66 def get_least_loaded_substation(): 1 usage
67     """Find the substation with the lowest current load"""
68     with load_lock:
69         if not substation_loads:
70             return SUBSTATIONS[0] # Default to first substation
71
72         min_load = float('inf')
73         best_substation = None
74
75         for substation in SUBSTATIONS:
76             load = substation_loads.get(substation['id'], 0)
77             if load < min_load:
78                 min_load = load
79                 best_substation = substation
80
81         return best_substation if best_substation else SUBSTATIONS[0]
```

rt-grid-load-balancer > load\_balancer > main.py

```
charge_request_service\main.py  test.py  load_balancer\main.py ×
66 def get_least_loaded_substation(): 1 usage
68     with load_lock:
69         min_load = load
70         best_substation = substation
71
72     return best_substation if best_substation else SUBSTATIONS[0]
73
74 @app.route('/route_charge', methods=['POST'])
75 def route_charge():
76     """Route charging request to the least loaded substation"""
77     try:
78         data = request.get_json()
79
80         if not data:
81             return jsonify({'error': 'No data provided'}), 400
82
83         # Find the best substation
84         best_substation = get_least_loaded_substation()
85
86         logger.info(f"Routing charge request to {best_substation['id']} "
87                   f"(load: {substation_loads.get(best_substation['id'], 0)})")
88
89         # Forward the request to the selected substation
90         response = requests.post(
91             url=f"{best_substation['url']}/charge",
92             json=data,
93             timeout=30
94         )
95
96         if response.status_code == 200:
97             result = response.json()
98             result['routed_by'] = 'load_balancer'
99             result['substation_id'] = best_substation['id']
100             result['substation_load_before'] = substation_loads.get(best_substation['id'], 0)
101             return jsonify(result), 200
102         else:
103             logger.error(f"Substation {best_substation['id']} returned error: {response.status_code}")
104             return jsonify({'error': 'Substation processing failed'}), 500
105
106     except requests.RequestException as e:
107         logger.error(f"Connection error to substation: {str(e)}")
108         return jsonify({'error': 'Substation unavailable'}), 503
109
110
111
112
113
114
115
116
117
mart-grid-load-balancer > load_balancer > main.py
```

## Load tester/test.py

```
import requests
import json
import time
import random
import threading
from datetime import datetime, timedelta
import sys

# Configuration
CHARGE_REQUEST_URL = "http://localhost:8000/charge"
LOAD_BALANCER_URL = "http://localhost:8080/status"
TOTAL_REQUESTS = 100
CONCURRENT_THREADS = 10
RUSH_HOUR_DURATION = 60 # seconds

# Statistics tracking
stats = {
    'total_requests': 0,
    'successful_requests': 0,
    'failed_requests': 0,
    'rejected_requests': 0,
    'response_times': [],
    'start_time': None,
    'end_time': None
}

stats_lock = threading.Lock()

def log_with_timestamp(message):
    """Print message with timestamp"""
    print(f"[{datetime.now().strftime('%H:%M:%S')}] {message}")

def generate_charging_request():
    """Generate a realistic EV charging request"""
    vehicle_id = f"EV_{random.randint(1000, 9999)}"

    # Realistic charge amounts (in kW)
    charge_amounts = [7, 11, 22, 50, 100, 150] # Common EV charging
    rates
    charge_amount = random.choice(charge_amounts)

    # Priority distribution (most requests are normal priority)
    priority_choices = ['low', 'normal', 'normal', 'normal', 'high']
    priority = random.choice(priority_choices)

    return {
        'vehicle_id': vehicle_id,
        'charge_amount': charge_amount,
        'priority': priority,
        'request_time': datetime.now().isoformat()
    }

def send_charge_request():
    """Send a single charge request and track statistics"""
    global stats
```

```

try:
    request_data = generate_charging_request()
    start_time = time.time()

    response = requests.post(
        CHARGE_REQUEST_URL,
        json=request_data,
        timeout=30
    )

    end_time = time.time()
    response_time = end_time - start_time

    with stats_lock:
        stats['total_requests'] += 1
        stats['response_times'].append(response_time)

        if response.status_code == 200:
            stats['successful_requests'] += 1
            result = response.json()
            log_with_timestamp(f"✓ Vehicle
{request_data['vehicle_id']} charged at "
                             f"{result.get('substation_id',
'unknown')}} "
                             f"({request_data['charge_amount']}kW,
{response_time:.2f}s)")
            elif response.status_code == 503:
                stats['rejected_requests'] += 1
                log_with_timestamp(f"△ Vehicle
{request_data['vehicle_id']} rejected - "
                                 f"insufficient capacity
({response_time:.2f}s)")
            else:
                stats['failed_requests'] += 1
                log_with_timestamp(f"X Vehicle
{request_data['vehicle_id']} failed - "
                                 f"HTTP {response.status_code}
({response_time:.2f}s)")

    except requests.RequestException as e:
        with stats_lock:
            stats['total_requests'] += 1
            stats['failed_requests'] += 1
        log_with_timestamp(f"X Request failed: {str(e)}")
    except Exception as e:
        with stats_lock:
            stats['total_requests'] += 1
            stats['failed_requests'] += 1
        log_with_timestamp(f"X Unexpected error: {str(e)}")

def worker_thread():
    """Worker thread that continuously sends requests during rush
hour"""
    while True:
        current_time = datetime.now()

```

```

        if stats['end_time'] and current_time >= stats['end_time']:
            break

        send_charge_request()

        # Random delay between requests (0.1 to 2 seconds)
        time.sleep(random.uniform(0.1, 2.0))

def monitor_system_status():
    """Monitor and log system status during load testing"""
    log_with_timestamp("Starting system monitoring...")

    while True:
        current_time = datetime.now()
        if stats['end_time'] and current_time >= stats['end_time']:
            break

        try:
            response = requests.get(Load_Balancer_URL, timeout=5)
            if response.status_code == 200:
                data = response.json()
                loads = data.get('substation_loads', {})
                load_info = ', '.join([f"{k}: {v}" for k, v in
loads.items()])
                log_with_timestamp(f"📊 Substation loads: {load_info}")
            else:
                log_with_timestamp(f"⚠️ Failed to get load balancer
status")
        except Exception as e:
            log_with_timestamp(f"⚠️ Error monitoring system: {str(e)}")

        time.sleep(10) # Monitor every 10 seconds

def print_final_statistics():
    """Print comprehensive test results"""
    print("\n" + "="*80)
    print("LOAD TEST RESULTS")
    print("="*80)

    duration = (stats['end_time'] -
stats['start_time']).total_seconds()

    print(f"Test Duration: {duration:.1f} seconds")
    print(f"Total Requests: {stats['total_requests']}")
    print(f"Successful Requests: {stats['successful_requests']}")
    print(f"({stats['successful_requests']/stats['total_requests']*100:.1f}%)")
    print(f"Rejected Requests: {stats['rejected_requests']}")
    print(f"({stats['rejected_requests']/stats['total_requests']*100:.1f}%)")
    print(f"Failed Requests: {stats['failed_requests']}")
    print(f"({stats['failed_requests']/stats['total_requests']*100:.1f}%)")
    print(f"Requests per Second: {stats['total_requests']/duration:.2f}")

    if stats['response_times']:
        avg_response_time = sum(stats['response_times']) /
len(stats['response_times'])

```



```

        min_response_time = min(stats['response_times'])
        max_response_time = max(stats['response_times'])

        print(f"\nResponse Times:")
        print(f"Average: {avg_response_time:.3f}s")
        print(f"Minimum: {min_response_time:.3f}s")
        print(f"Maximum: {max_response_time:.3f}s")

    print("\n" + "="*80)

def simulate_rush_hour():
    """Simulate a rush hour of EV charging requests"""
    print("🚗 Smart Grid Load Balancer - Rush Hour Simulation")
    print("="*60)

    # Initialize timing
    stats['start_time'] = datetime.now()
    stats['end_time'] = stats['start_time'] +
timedelta(seconds=RUSH_HOUR_DURATION)

    log_with_timestamp(f"Starting {RUSH_HOUR_DURATION}s rush hour
simulation with {CONCURRENT_THREADS} concurrent threads")

    # Start monitoring thread
    monitor_thread = threading.Thread(target=monitor_system_status,
daemon=True)
    monitor_thread.start()

    # Start worker threads
    threads = []
    for i in range(CONCURRENT_THREADS):
        thread = threading.Thread(target=worker_thread, daemon=True)
        thread.start()
        threads.append(thread)
        log_with_timestamp(f"Started worker thread {i+1}")

    # Wait for all threads to complete
    for thread in threads:
        thread.join()

    log_with_timestamp("Rush hour simulation completed!")

    # Print final statistics
    print_final_statistics()

def run_simple_test():
    """Run a simple test with a few requests"""
    print("🔧 Running simple connectivity test...")

    for i in range(5):
        log_with_timestamp(f"Sending test request {i+1}/5")
        request_data = generate_charging_request()

        try:
            response = requests.post(CHARGE_REQUEST_URL,
json=request_data, timeout=10)

```

```

        if response.status_code == 200:
            result = response.json()
            log_with_timestamp(f"✓ Success: {result.get('substation_id', 'unknown')}")
        else:
            log_with_timestamp(f"△ HTTP {response.status_code}: {response.text}")
        except Exception as e:
            log_with_timestamp(f"✗ Error: {str(e)}")

        time.sleep(2)

if __name__ == "__main__":
    if len(sys.argv) > 1 and sys.argv[1] == "simple":
        run_simple_test()
    else:
        simulate_rush_hour()

```

### substation service/main.py

```

from flask import Flask, request, jsonify
import logging
import os
import time
import threading
import random
from datetime import datetime, timedelta

app = Flask(__name__)
logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

# Configuration
SUBSTATION_ID = os.getenv('SUBSTATION_ID', 'substation_unknown')
MAX_CAPACITY = int(os.getenv('MAX_CAPACITY', '100')) # Maximum
charging capacity
CHARGE_PROCESSING_TIME = int(os.getenv('CHARGE_PROCESSING_TIME', '10'))
# seconds

# Current load tracking
current_load = 0
charging_sessions = {} # Store active charging sessions
load_lock = threading.Lock()

def simulate_charging_completion():
    """Background thread to simulate charging completion"""
    global current_load, charging_sessions

    while True:
        try:
            current_time = datetime.now()
            completed_sessions = []

            with load_lock:
                for session_id, session_data in

```

```

charging_sessions.items():
    if current_time >= session_data['end_time']:
        completed_sessions.append(session_id)
        current_load -= session_data['charge_amount']
        logger.info(f"Charging completed for session
{session_id}, "
                    f"load reduced by
{session_data['charge_amount']}")

    # Remove completed sessions
    for session_id in completed_sessions:
        del charging_sessions[session_id]

    # Ensure load doesn't go negative
    if current_load < 0:
        current_load = 0

    time.sleep(2)  # Check every 2 seconds

except Exception as e:
    logger.error(f"Error in charging completion thread:
{str(e)}")
    time.sleep(5)

@app.route('/charge', methods=['POST'])
def process_charge():
    """Process a charging request"""
    global current_load, charging_sessions

    try:
        data = request.get_json()

        if not data:
            return jsonify({'error': 'No data provided'}), 400

        # Validate required fields
        required_fields = ['vehicle_id', 'charge_amount', 'priority']
        for field in required_fields:
            if field not in data:
                return jsonify({'error': f'Missing required field:
{field}'}), 400

        charge_amount = float(data['charge_amount'])
        vehicle_id = data['vehicle_id']
        priority = data.get('priority', 'normal')

        with load_lock:
            # Check if we can handle this charge request
            if current_load + charge_amount > MAX_CAPACITY:
                logger.warning(f"Charge request rejected - would exceed
capacity "
                               f"(current: {current_load}, requested:
{charge_amount}, max: {MAX_CAPACITY})")
                return jsonify({
                    'error': 'Insufficient capacity',
                    'current_load': current_load,
                    'max_capacity': MAX_CAPACITY,

```

```

        'available_capacity': MAX_CAPACITY - current_load
    }), 503

    # Accept the charging request
    session_id =
f"{SUBSTATION_ID}_{vehicle_id}_{int(time.time())}"

    # Calculate charging duration based on amount and priority
    base_duration = CHARGE_PROCESSING_TIME
    if priority == 'high':
        duration = max(base_duration * 0.7, 5) # 30% faster
for high priority
    elif priority == 'low':
        duration = base_duration * 1.3 # 30% slower for low
priority
    else:
        duration = base_duration

    # Add some randomness to simulate real-world variation
    duration = duration * (0.8 + random.random() * 0.4) # ±20%
variation

    start_time = datetime.now()
    end_time = start_time + timedelta(seconds=duration)

    # Update load and add session
    current_load += charge_amount
    charging_sessions[session_id] = {
        'vehicle_id': vehicle_id,
        'charge_amount': charge_amount,
        'priority': priority,
        'start_time': start_time,
        'end_time': end_time,
        'duration': duration
    }

    logger.info(f"Started charging session {session_id} for
vehicle {vehicle_id}, "
               f"amount: {charge_amount}, duration:
{duration:.1f}s, "
               f"new load: {current_load}")

    return jsonify({
        'status': 'accepted',
        'session_id': session_id,
        'substation_id': SUBSTATION_ID,
        'vehicle_id': vehicle_id,
        'charge_amount': charge_amount,
        'estimated_duration': duration,
        'start_time': start_time.isoformat(),
        'current_load': current_load,
        'max_capacity': MAX_CAPACITY
    }), 200

except ValueError as e:
    return jsonify({'error': f'Invalid charge amount: {str(e)}'}),
400

```

```

        except Exception as e:
            logger.error(f"Unexpected error in charge processing: {str(e)}")
            return jsonify({'error': 'Internal server error'}), 500

@app.route('/metrics', methods=['GET'])
def metrics():
    """Expose metrics in Prometheus format"""
    with load_lock:
        metrics_text = f"# HELP substation_current_load Current charging load of the substation\n"
        metrics_text += f"# TYPE substation_current_load gauge\n"
        metrics_text += f"substation_current_load {current_load}\n"

        metrics_text += f"# HELP substation_max_capacity Maximum capacity of the substation\n"
        metrics_text += f"# TYPE substation_max_capacity gauge\n"
        metrics_text += f"substation_max_capacity {MAX_CAPACITY}\n"

        metrics_text += f"# HELP substation_active_sessions Number of active charging sessions\n"
        metrics_text += f"# TYPE substation_active_sessions gauge\n"
        metrics_text += f"substation_active_sessions {len(charging_sessions)}\n"

        utilization = (current_load / MAX_CAPACITY) * 100 if MAX_CAPACITY > 0 else 0
        metrics_text += f"# HELP substation_utilization_percent Capacity utilization percentage\n"
        metrics_text += f"# TYPE substation_utilization_percent gauge\n"
        metrics_text += f"substation_utilization_percent {utilization:.2f}\n"

    return metrics_text, 200, {'Content-Type': 'text/plain'}

@app.route('/health', methods=['GET'])
def health_check():
    """Health check endpoint"""
    return jsonify({'status': 'healthy', 'substation_id': SUBSTATION_ID}), 200

@app.route('/status', methods=['GET'])
def status():
    """Get detailed status of the substation"""
    with load_lock:
        return jsonify({
            'substation_id': SUBSTATION_ID,
            'current_load': current_load,
            'max_capacity': MAX_CAPACITY,
            'utilization_percent': (current_load / MAX_CAPACITY) * 100 if MAX_CAPACITY > 0 else 0,
            'active_sessions': len(charging_sessions),
            'available_capacity': MAX_CAPACITY - current_load,
            'sessions': {k: {
                'vehicle_id': v['vehicle_id'],
                'charge_amount': v['charge_amount'],
            }} for k, v in charging_sessions.items()
        })

```

```

                'priority': v['priority'],
                'remaining_time': (v['end_time'] -
datetime.now()).total_seconds()
            } for k, v in charging_sessions.items()},
            'timestamp': datetime.now().isoformat()
        )), 200

if __name__ == '__main__':
    # Start the charging completion simulation thread
    completion_thread =
threading.Thread(target=simulate_charging_completion, daemon=True)
    completion_thread.start()

    logger.info(f"Starting substation {SUBSTATION_ID} with capacity
{MAX_CAPACITY}")
    app.run(host='0.0.0.0', port=8001, debug=False)

```

## docker-compose.yml

```

services:
  # Charge Request Service - Public entry point
  charge_request_service:
    build:
      context: ./charge_request_service
      dockerfile: Dockerfile
    ports:
      - "8000:8000"
    environment:
      - LOAD_BALANCER_URL=http://load_balancer:8080
    depends_on:
      - load_balancer
    networks:
      - smart-grid
    restart: unless-stopped

  # Load Balancer - Core routing logic
  load_balancer:
    build:
      context: ./load_balancer
      dockerfile: Dockerfile
    ports:
      - "8080:8080"
    depends_on:
      - substation_1
      - substation_2
      - substation_3
    networks:
      - smart-grid
    restart: unless-stopped

  # Substation Services - Multiple replicas with different capacities
  substation_1:

```

```
    build:
      context: ./substation_service
      dockerfile: Dockerfile
    environment:
      - SUBSTATION_ID=substation_1
      - MAX_CAPACITY=80
      - CHARGE_PROCESSING_TIME=8
    expose:
      - "8001"
    networks:
      - smart-grid
    restart: unless-stopped

substation_2:
  build:
    context: ./substation_service
    dockerfile: Dockerfile
  environment:
    - SUBSTATION_ID=substation_2
    - MAX_CAPACITY=120
    - CHARGE_PROCESSING_TIME=10
  expose:
    - "8001"
  networks:
    - smart-grid
  restart: unless-stopped

substation_3:
  build:
    context: ./substation_service
    dockerfile: Dockerfile
  environment:
    - SUBSTATION_ID=substation_3
    - MAX_CAPACITY=100
    - CHARGE_PROCESSING_TIME=12
  expose:
    - "8001"
  networks:
    - smart-grid
  restart: unless-stopped

# Prometheus - Metrics collection
prometheus:
  image: prom/prometheus:latest
  ports:
    - "9090:9090"
  volumes:
    -
./monitoring/prometheus/prometheus.yml:/etc/prometheus/prometheus.yml
  - prometheus_data:/prometheus
  command:
    - '--config.file=/etc/prometheus/prometheus.yml'
    - '--storage.tsdb.path=/prometheus'
    - '--web.console.libraries=/etc/prometheus/console_libraries'
    - '--web.console.templates=/etc/prometheus/consoles'
    - '--storage.tsdb.retention.time=200h'
    - '--web.enable-lifecycle'
```

```

    networks:
      - smart-grid
    restart: unless-stopped

# Grafana - Metrics visualization
grafana:
  image: grafana/grafana:latest
  ports:
    - "3000:3000"
  volumes:
    - grafana_data:/var/lib/grafana
    - ./monitoring/grafana:/etc/grafana/provisioning
  environment:
    - GF_SECURITY_ADMIN_USER=admin
    - GF_SECURITY_ADMIN_PASSWORD=admin
    - GF_USERS_ALLOW_SIGN_UP=false
  networks:
    - smart-grid
  depends_on:
    - prometheus
  restart: unless-stopped

networks:
  smart-grid:
    driver: bridge

volumes:
  prometheus_data:
  grafana_data:

```

## prometheus.yml

```

global:
  scrape_interval: 15s
  evaluation_interval: 15s

rule_files:
  # - "first_rules.yml"
  # - "second_rules.yml"

scrape_configs:
  # The job name is added as a label `job=<job_name>` to any timeseries
  # scraped from this config.
  - job_name: 'prometheus'
    static_configs:
      - targets: ['localhost:9090']

  # Scrape substation metrics
  - job_name: 'substations'
    static_configs:
      - targets:
          - 'substation_1:8001'
          - 'substation_2:8001'
          - 'substation_3:8001'
    scrape_interval: 5s

```



```
    metrics_path: '/metrics'

# Scrape load balancer metrics
- job_name: 'load_balancer'
  static_configs:
    - targets: ['load_balancer:8080']
  scrape_interval: 5s
  metrics_path: '/metrics'

# Optional: Monitor the charge request service
- job_name: 'charge_request_service'
  static_configs:
    - targets: ['charge_request_service:8000']
  scrape_interval: 10s
  metrics_path: '/health'
```